Measuring Model Substitution in LLM API Resellers: A 14-Day Study of 171 Endpoints

Bazaarlink Research[1]

[1]bazaarlink.ai — corresponding repository:, github.com/Bazaarlinkorg/LLMprobe-engine

Abstract

The market for resold access to large-language-model (LLM) APIs has expanded rapidly since 2024, driven by upstream providers' regional service-availability policies. Because upstream providers cannot authenticate which back-end model a third-party reseller actually routes to, and because protocol-layer responses are nearly indistinguishable across models, downstream buyers cannot directly detect silent model substitution. We design a four-layer black-box detection system based on independent behavioral fingerprints: a surface family classifier, a behavior column with self-claim zeroed, a sub-model deterministic match (cutoff / capability / refusal-prefix), and a behavioral-vector extension covering refusal-boundary ladder, formatting idiosyncrasy, and calibrated uncertainty. We apply it to 171 distinct reseller endpoints over a 14-day window (2026-04-12 to 2026-04-25; 625 probe runs). In a fixed-size synthetic stress test the system achieves 94.4% intra-family true-positive rate (n=18; Wilson 95% CI [74.2%, 99.0%]) at 0% false-positive rate (n=6). In the field, **9.9% of probed endpoints (17/171) show at least one detected substitution event; under a stricter criterion ($n \geq 5$ probes, $\geq$20% violation rate) 1.3% (2/149) meet our persistent-substitution criterion**. The central empirical contribution is the case analysis: R1 exhibits dynamic routing across multiple back ends under one advertised model, R2 exhibits systematic version-tag generalization, and R6 shows that single-session incident-style substitution can still be captured by an internally consistent evidence chain. Observed substitutions sort into five categories: cross-family impersonation, silent downgrade, silent upgrade, version-tag forgery, and provider-injected behavioral constraint. We complement Liu et al. (arXiv:2604.08407) [5] — who measure malicious-intermediary security attacks on the LLM supply chain — by quantifying the orthogonal axis of model-substitution quality fraud, and close with three disclosure-duty proposals.

Keywords: LLM API resale; model substitution; behavioral fingerprinting; consumer protection; black-box auditing.

1 Introduction

1.1 A market born of access restrictions

Anthropic and OpenAI publicly disclose, in their service terms and acceptable-use policies [2, 7], that their flagship products are unavailable for direct purchase in certain countries and regions. Demand for these tools has not abated in those regions; the productivity arms race between developers in restricted regions and their global peers has, if anything, intensified the demand. The supply–demand gap so created has, over the past twelve months, given rise to a high-growth resale industry: third-party operators stand up OpenAI-compatible HTTP endpoints in accessible jurisdictions, back the endpoints with first-party API keys, and resell access at a markup to downstream users in restricted regions. Such resale endpoints typically advertise support for flagship models (Claude Opus 4.7, GPT-5.4, Gemini 3.1 Pro), and present an OpenAI-compatible interface so that downstream code originally targeted at first-party SDKs may switch over without modification.

This is not the familiar pattern of a sanctioned API reseller. Sanctioned resellers (e.g., on cloud-marketplaces) typically have formal contracts with the upstream provider, contractual SLAs, and an interface that surfaces the upstream provenance. The resellers studied in this paper instead operate as *unauthorized gray-market intermediaries*: upstream providers have not authorized resale, the resellers' upstream API keys cannot be reconciled to downstream buyers in the upstream's books, and downstream users routinely *assume* service homogeneity from interface compatibility.

1.2 The collapse of trust assumptions

Downstream buyers' trust in such intermediaries traditionally rests on four assumptions: (1) **in-**

terface: the endpoint implements the OpenAI-compatible protocol, therefore its behavior approximates first-party; (2) **self-claim**: a request that passes `model: claude-opus-4-7` and receives a response self-identifying as Claude indicates the back end truly is that model; (3) **branding**: a website displaying upstream logos and customer service claiming "official partnership" sources from the official upstream; (4) **tier consistency**: resellers will not advertise a high-tier model while routing to a low-tier model. We will demonstrate that all four assumptions fail in the present resale market.

1.3 Research questions

RQ1 (Prevalence) What fraction of currently active LLM API reseller endpoints exhibit externally observable model substitution? **RQ2 (Typology)** Can the observed substitution behaviors be organized into reproducible categories? **RQ3 (Detection feasibility)** Under purely black-box observation, can a detection method achieve simultaneously high true-positive rate and low false-positive rate? **RQ4 (Consumer impact)** What concrete harms does substitution impose on downstream buyers, and what are the implications for regulators and upstream providers?

1.4 Contributions and relation to prior work

The contributions are: (1) a first systematic measurement of the resale market for model-substitution behavior; (2) a reproducible black-box detection methodology achieving 94.4% intra-family TP at 0% FP under fixed-size synthetic tests; (3) a five-category taxonomy of substitution behaviors; (4) a consumer-protection policy analysis with three immediately implementable disclosure-duty proposals.

Relation to prior work. The closest contemporary measurement is Liu et al. [5], who measure 428 LLM API routers for *active security attacks*: payload injection, secret/credential exfiltration, evasion tactics, and crypto-asset access. Our work measures the orthogonal axis: routers that perform *passive quality substitution* — silently routing an advertised premium model to a lower-tier or different one — without active intrusion on the user's system. Together, the two studies span the principal threat surfaces of the LLM API supply chain: Liu et al. on what the router does *to the user*, and the present paper on what the router does *with the model*.

For the technical building blocks of layer ④ (Section 3.6) we draw on three threads. The refusal-boundary ladder is informed by mechanistic work on refusal direction in safety-tuned LLMs [3]. The formatting fingerprint follows the observation that stylistic preferences are pre-training-baked and resist runtime instruction [6]. The calibrated-uncertainty channel applies the round-rate intuition from Kadavath et al.'s study of LLM self-knowledge [4]. The deterministic sub-model matcher of layer ③ extends behavioral fingerprinting along the lines of LLMmap [8] but adapts it to within-family discrimination rather than across-family identification. The economic framing in Section 7 builds on Akerlof's seminal analysis of asymmetric information in markets where quality is unobservable to buyers [1]. Methodological design draws on Paxson's framework for sound Internet measurement [9].

1.5 Structure

Section 2 lays out the threat model. Section 3 describes the four-layer detection method. Section 4 outlines the engineering implementation. Section 5 reports synthetic and field measurement results. Section 6 organizes substitution behaviors into a taxonomy with anchor cases including three deep case studies (R1, R2, R6). Section 7 analyzes consumer harm and policy implications. Section 8 discusses limitations. Section 9 concludes.

2 Threat Model

2.1 Adversary types

We classify substitution-capable adversaries (resellers) by commercial motivation:

- **Arbitrage adversary.** Advertises model M_{high}; routes to a lower-tier same-family model M_{low} (e.g., advertises Opus 4.7, routes to Sonnet 4.6). Motivation: margin expansion through quality-tier substitution.
- **Silent-downgrade adversary.** Advertises M_{high} and initially routes to M_{high}, but switches to M_{low} under operational pressure (upstream-quota exhaustion, routing fallback, or vendor policy change).
- **Cross-family adversary.** Advertises a model in family F_a (Anthropic Claude); routes to a model in family F_b (OpenAI GPT, Google Gemini, Z-AI GLM). Motivation: cross-family arbitrage, or fail-over when F_a accounts are blocked.

A non-pecuniary fourth pattern is observed:

- **Behavior-injection adversary.** Injects a system prompt at the upstream layer (e.g., Amazon Q Developer / Kiro coding-agent wrapper) that constrains an otherwise general-purpose model to a narrow domain. Downstream users do receive the advertised model, but its behavioral freedom is silently restricted.

2.2 Adversary capabilities

The adversary has full control over: request routing (upstream key + model parameter), request rewriting (system prompt injection), response rewriting, header injection, and streaming forgery. The adversary does *not* possess the ability to rewrite the upstream model's internal reasoning (logprobs, token-count physics) or to forge the full behavioral fingerprint of an upstream model in absence of that model's participation. The adversary may rewrite *what the response says*, but cannot readily fabricate a coherent set of responses spanning multiple independent fingerprint channels.

2.3 Attack-surface classification

Mapping the substitution to its injection point in the request lifecycle:

3 Methodology: Four-Layer Independent Fingerprint Detection

3.1 Design principles

P1 Multi-layer independence. Each detection layer uses an independent signal channel; an adversary must simultaneously forge consistent behavior across channels to escape detection. **P2 Black-box accessibility.** All channels are obtained via downstream HTTP API alone; no privileged access is required. **P3 Quantified confidence.** The system outputs not only a Boolean verdict but a quantified confidence band and per-layer evidence trail.

3.2 Architecture overview

3.3 Layer ①: Surface fingerprint

Principle. The overall textual response (including any self-identifying language it contains) carries family-discriminative information with a small number of probes.

Implementation. A linear classifier maps response features to family weights (Anthropic / OpenAI / Google / xAI / other) with weights pretrained on known-model traffic.

Strength. Computationally cheap; correct in the vast majority of unattacked cases. **Weakness.** Trivially polluted by self-claim. An adversary need only inject "you are Claude Opus 4.7" into the upstream system prompt to drive surface output to Anthropic = 100%. **Layer ① must not be used in isolation.**

3.4 Layer ②: Behavior column (self-claim zeroed)

Principle. Stripping self-identifying short phrases from the response and re-classifying yields a column that is largely robust against prompt-layer self-claim forgery. Adversaries rewrite self-identification far more often than they rewrite style features.

Implementation. Two parallel classifiers: one over the raw response, one after regex-deletion of a curated self-identification phrase set.

Strength. Robust against prompt-layer attacks within its band of confidence. **Weakness.** Powerless against pure routing-layer attacks: when the back end *is* a different family, both columns will agree and indicate the wrong family relative to the user's claim. Such cases are caught at the verdict layer by comparing claim against behavior-column truth.

3.5 Layer ③: Sub-model deterministic match

Principle. Within a family, sibling sub-models (e.g., Claude Opus 4.5 / 4.6 / 4.7) differ measurably along three observable dimensions: training-data cutoff, capability boundary, and refusal-prefix style.

Implementation. A baseline database stores, per target model, its measurements along: *cutoff* (response to "what is your training-data cutoff date?"), *capability boundary* (a fixed set of factual questions including at least one whose answer changes across plausible cutoff dates), and *refusal prefix* (observed leading text and length when the model refuses a standardized sensitive request). A heuristic match score selects the best-matching sub-model when score ≥ 0.60 and gap to runner-up ≥ 0.05; otherwise abstain.

Strength. Strong against pure routing-layer attacks: a different family scores zero across the board; a different sub-model differs on at least one dimension. **Weakness.** Each dimension can be precisely targeted by a prompt-layer attack ("your cutoff is 2025-04"); the layer must be paired with layer ④ for robustness.

3.6 Layer ④: Behavioral-vector extension

Principle. Layer ③ is augmented with channels that are harder to forge by system-prompt manipulation. Three vectors:

- **Refusal ladder** (8 prompts of monotonically increasing sensitivity, from compliant to high-sensitivity, with CBRN material replaced by cybercrime equivalents). Each rung is classified $\{\text{comply}, \text{partial}, \text{refuse}\}$; the resulting 8-dim vector empirically discriminates Claude Opus

Surface	Injection point	Forgery ceiling
Routing layer	upstream key / `model`	full back-end replacement
Prompt layer	upstream system prompt	forge self-claim; not refusal floor
Response layer	post-process upstream response	rewrite text; not metadata

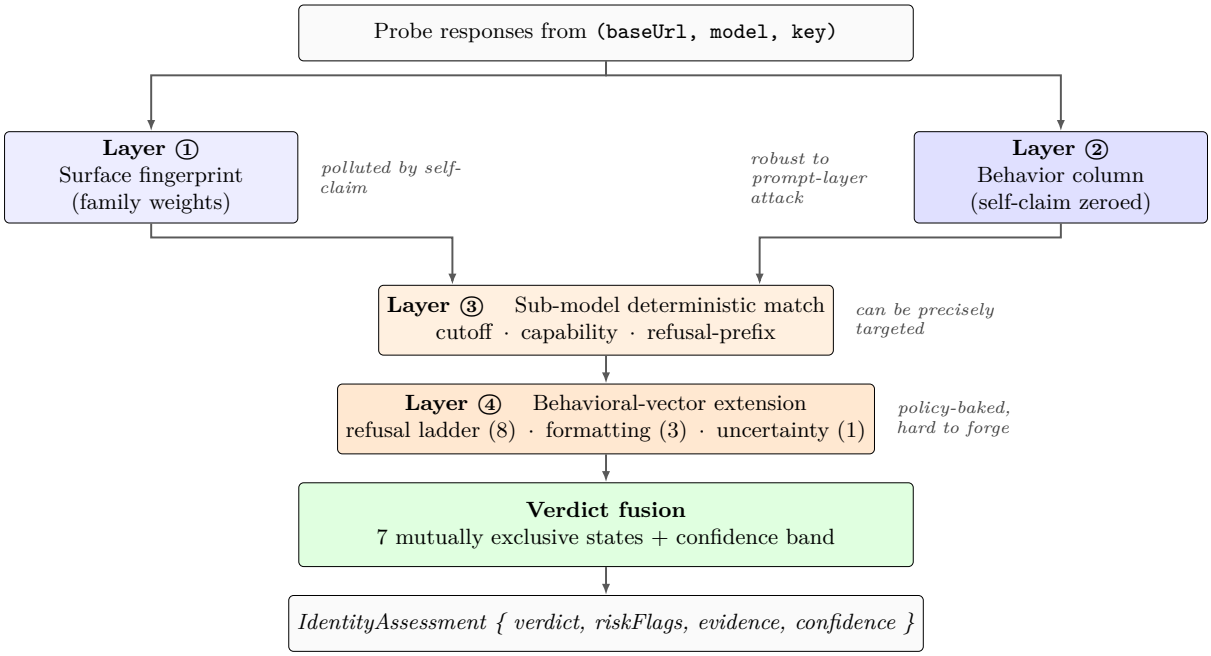

Figure 1: The four-layer independent-fingerprint detection architecture (§3). Each layer uses an independent signal channel; an adversary must simultaneously forge consistent behavior across multiple channels to escape detection. The verdict fusion emits one of seven mutually exclusive states with a confidence band.

4.7 from Anthropic siblings at rung L6 (involving methods of self-harm).

- **Formatting idiosyncrasy** (3 probes): preferred bullet glyph (`-` / `*` / `1.`), markdown header depth, code-block language tag.
- **Calibrated uncertainty** (1 probe): when asked to express its confidence numerically, both the value (e.g., 65%) and whether the value is a round multiple of ten. The latter discriminates same-family siblings such as GPT-5.5 (isRoundRate=1.0) from GPT-5.3-codex (isRoundRate=0.33).

Strength. Refusal boundaries are determined by the underlying safety policy and resist system-prompt rewriting; formatting preferences derive from training distribution rather than runtime instruction; numeric-rounding habits are highly robust against "pretend to be confident" prompt attacks. **Weakness.** Twelve additional probes are required.

3.7 Verdict fusion and seven-state output

4 Engineering Implementation

The methodology is implemented as an open-source TypeScript library, `@bazaarlink/probe-engine`, released under AGPL-3.0 (current version v0.7.1; see Section C). The library exposes the four classifiers as pure functions over recorded HTTP responses, plus a default probe suite of fixed prompts. A reproducer in possession of the open-source package, an OpenRouter (or equivalent) API key for baseline construction, and the published baseline snapshot can independently re-run the evaluation reported in Section 5; no centralized infrastructure is required.

For the field-measurement reported in Section 5.3, the authors operate a thin Web frontend over the library at `https://bazaarlink.ai/probe`. The frontend accepts a `(baseUrl, modelId, apiKey)` triple from a submitting user, executes the probe suite against the supplied endpoint, persists results to a `probe_history` table, and returns the verdict to the submitter. **This frontend is a measurement instrument used to collect the field data; it is not part of the methodology and is not**

Verdict	Trigger (simplified)
`clean_match`	Layers 1, 2, 3 (or 4) all consistent with claim, conf. ≥ 0.85
`match`	Layers 1, 2 consistent; layer 3 (or 4) abstain or unavailable
`clean_match_submodel_mismatch`	Family consistent but sub-model layer disagrees
`plain_mismatch`	Behavior-column family $\neq$ claimed family
`mismatch`	Layer 3/4 disagrees with claim at high confidence
`spoof_selfclaim_forged`	Self-claim present, behavior-column family differs
`spoof_behavior_induced`	Style induced toward claim, layers 3/4 still detect back-end
`ambiguous/uncertain/insufficient_data`	Insufficient or contradictory signal

required for reproduction.

5 Quantitative Results

5.1 Baseline coverage

As of 2026-04-26 the sub-model baseline database contains 27 models across 6 families. Layer ③ coverage is complete for all 27 models through cutoff, capability, and refusal-prefix features. Layer ④ contains 22 full extended-behavior baselines, covering every model family that appeared as a downstream claim during the measurement window:

Family	Models	Layer 4 ready	Coverage
Anthropic	6	6	100%
OpenAI	10	10	100%
Google	5	5	100%
Total	**27**	**22**	**81%**

The five models without layer ④ snapshots are retained in the layer ③ database but excluded from layer ④ analysis because they had no downstream claims in the field-measurement corpus. This keeps the denominator transparent: the public snapshot supports complete reproduction of all layer ④ verdicts reported in Section 5, while preserving layer ③ baselines for future measurements when additional families appear in the market. Baseline rows are versioned by model identifier and snapshot date; the open-source release includes a file-based loader so reproduction does not depend on the authors' SaaS database.

5.2 Synthetic stress test

The synthetic evaluation executed 18 intra-family substitution rounds and 6 genuine-model stress rounds. We honestly disclose: n=18 yields a 95% Wilson confidence interval of $[74.2\%, 99.0\%]$ — sample size is **insufficient to claim a tight point estimate**.

The single missed case was Anthropic Claude Opus 4.6 under injection-mode attack pushing its L6 refusal boundary to Opus 4.7's level. This is the layer ④ injection ceiling, addressed in Section 8.2.

5.3 14-day field deployment

Volume. Total probe runs: 625. Distinct `baseUrl` endpoints: 171. Distinct claimed `modelId` values: 47. Admin-verified-clean stamps: 6.

Verdict distribution (n=561 with verdict). Strict reading: **12 cases (2.1%) are high-confidence spoof verdicts** (the two `spoof_*` states combined). Broader reading including `plain_mismatch`, `mismatch`, and `clean_match_submodel_mismatch`: **100 cases (17.8%) exhibit some form of model deviation**. The two spoof states split evenly (6+6), consistent with the dual attack-surface hypothesis.

Endpoint-level violation distribution. Treating each `baseUrl` as a single endpoint and bucketing by 14-day violation rate:

Principal finding (RQ1). Among the 149 higher-evidence endpoints ($n \geq 5$), **7** ($\approx 4.7\%$) show $\geq$1 violation; **2** ($\approx 1.3\%$) show $\geq$20% sustained violation rate. Relaxing to $n \geq 1$, $17/171 \approx 9.9\%$ show $\geq$1 violation, but at this relaxation the impact of one-off false positives cannot be ruled out.

6 A Taxonomy of Substitution Behavior

Five categories suffice to organize observed behavior. Each is illustrated by an anchor case from the field; Section 6.7 provides deep case studies for R1, R2, and R6.

6.1 Cross-family impersonation

Advertise model in F_a; route to F_b. **Anchor (R6).** A probe of `claude-opus-4-6` returned a response whose behavior-column family weights placed `openai` = 0.96; sub-model deterministic match identified GPT-5.4. Verdict: `spoof_behavior_induced`, high confidence.

6.2 Intra-family silent downgrade

Advertise higher-tier; route to lower-tier same-family. **Anchor (R8).** Multiple probes of

Condition	n	TP / TN rate
Intra-family — naive system prompt	6	6/6 = 100%
Intra-family — forceful system prompt	6	6/6 = 100%
Intra-family — injection-mode system prompt	6	5/6 = 83%
Intra-family aggregate	**18**	**17/18 = 94.4%**
Genuine model under spoof prompt — flagged as spoof	6	**0/6 = 0%**

`verdict.status`	n	share
`clean_match`	279	49.7%
`match`	162	28.9%
`clean_match_submodel_mismatch`	35	6.2%
`plain_mismatch`	35	6.2%
`mismatch`	14	2.5%
`insufficient_data`	11	2.0%
`uncertain`	7	1.2%
`spoof_selfclaim_forged`	6	1.1%
`spoof_behavior_induced`	6	1.1%
`ambiguous`	6	1.1%

`claude-opus-4-7` stably identified by layer ④ as Claude Opus 4.6 at score ≥ 0.84. L6 refusal pattern and formatting preferences match 4.6.

6.3 Intra-family silent upgrade

Advertise lower-tier same-family; route to higher-tier. **Research observation.** Upgrade-direction substitution defies the obvious commercial incentive. Two non-fraudulent explanations are plausible: (a) the reseller's upstream API key is provisioned only for the higher-tier model; (b) operational error. **Anchor (R2).** Multiple probes of `claude-opus-4-6` identified the back end as Claude Opus 4.7 at 0.89.

6.4 Version-tag forgery

Advertise dated or preview-tagged version; route to same-family neighboring version. The harm: enterprises pin model versions through dated tags to lock down behavior across regression tests; tag forgery causes such enterprises to believe they are pinned while in fact receiving a different training snapshot. **Anchor (R2/R3).**

6.5 Provider behavior injection

Inject system prompt at upstream that constrains an otherwise general-purpose model into a narrow domain. The user does receive M, but its behavioral freedom is silently restricted. **Anchor.** A "KIRO (Amazon Q Developer) detection" flag activated 23 times across multiple endpoints during the 14-day window.

6.6 Endpoint summary

6.7 Deep case studies

R1 (`api.1xx.ai`) — dynamic-routing cross-family impersonation. R1 leads the measurement window in both volume (73 probes) and violation rate (61.6%, 45/73), and is the highest-evidence persistent-substitution candidate in our data. Two features distinguish R1's behavior: (a) a single advertised model maps to multiple distinct back-end models across different probes; (b) both prompt-layer and routing-layer attack surfaces are simultaneously observed at this endpoint.

Concretely, of five probes claiming `claude-opus-4-6`, two were identified as the genuine Claude Opus 4.6 (`clean_match`), one as GPT-5.4, one as GPT-5.4 Mini, and one as a non-Anthropic model whose specific identity could not be locked down. This distribution is inconsistent with any single static routing map; it is consistent with **dynamic load distribution**: the endpoint dispatches the same advertised claim to whichever back end is currently available.

R1 also exhibits the most complete prompt-layer forgery observed in our data. One probe of `claude-haiku-4-5-20251001` was classified `spoof_behavior_induced`: behavior-column identified GPT-5.3 Codex while the response style was being induced toward Claude. A separate probe of `gpt-5.3-codex-high` was classified `spoof_selfclaim_forged`: the response contained an explicit "I am GPT" self-identification that was contradicted by behavior-column family weights. **Both attack surfaces co-occurring at one endpoint suggest the operator possesses the engineering capacity to bypass deterministic sub-model matching (layer ③), yet remains caught by the behavioral-vector extension (layer ④).**

Operationally, R1 is important because it is not merely a mislabeled static proxy. A static proxy would consistently map a claimed model to a single neighboring back end, producing a stable mismatch pattern. R1 instead produces a many-to-one and one-to-many mapping: the same advertised model sometimes resolves to the genuine model and sometimes to multiple OpenAI-family alternatives. This is consistent with a routing layer that treats the `model` field as a soft label rather than a binding contract. For downstream users, the failure mode

Bucket	Endpoints	Cumulative probes
0% (clean, $n \geq 1$)	142	306
$< 20\%$ ($n \geq 5$)	5	106
20–50% ($n \geq 5$)	1	57
$\geq 50\%$ ($n \geq 5$)	1	73
$n < 5$, ≥ 1 violation (low evidence)	9	19
Other	13	64

Code	Masked domain	n	Viol.	Pattern
R1	`api.1xx.ai`	73	45	mixed
R2	`hbxxx.ai`	57	20	upgrade + version forgery
R3	`yunxx.ai`	27	5	version drift
R4	`api.linxxxxxai.com`	26	4	sporadic
R5	`api.gexxxxxxx.top`	7	1	dual-protocol
R6	`api.axxx.com`	6	1	cross-family
R7	`api.sanxxxxxai.top`	4	2	cross-family
R8	`api.aixxxxxx.com`	4	1	downgrade
R9	`us.noxxxxxx.com`	4	1	label forgery
R10–R15	(low evidence)	1–2	1–2	various

Table 1: Endpoints with detected violations during the 14-day window. Codename order reflects evidence strength. R10–R15 are placed in a low-evidence cohort and not counted toward the principal point estimates.

is intermittent and hard to reproduce: a complaint may be answered by a clean later call, even though the previous call was routed elsewhere.

R2 (`hbxxx.ai`) — version-tag generalization routing. R2 accumulated 57 probes and 20 violations (35.1%) in the 14-day window. Its violation profile differs from R1's: R2's back-end models are mostly genuine within their advertised families (Anthropic / Google), but the endpoint exhibits a systematic inconsistency in version-tag handling. We name this pattern **version-tag generalization routing**: the endpoint accepts arbitrary sub-model version tags — including discontinued versions, preview tags, and date-pinned versions — and uniformly routes them to the "most recent available version" within the same family.

Six probes of `claude-opus-4-5-20251101` were classified `clean_match_submodel_mismatch`, with behavior column identifying Claude Opus 4.5 (no date tag); 3 probes of `claude-sonnet-4-5-20250929` were uniformly identified as plain Claude Sonnet 4.5; probes of `gemini-3-pro-preview` were identified as Gemini 3.1 Pro. Most strikingly, 3 probes of `claude-sonnet-4-20250514` (a tag with no corresponding entry in Anthropic's actual product line) were all routed to Claude Haiku 4.5 — a cross-tier substitution that fits neither downgrade nor upgrade typology cleanly.

R2's pattern particularly harms enterprise users who pin behavior to dated version tags for regression-test stability. Should Anthropic ever update the underlying weights for Opus 4.5, R2's downstream users would silently receive a new training snapshot, invalidating their evaluation criteria without warning.

R2 is therefore not best understood as a simple "fake model" endpoint. Its dominant behavior is a compatibility illusion: the endpoint accepts model identifiers that look precise enough for production pinning, but internally collapses them onto a smaller set of available models. This is difficult for users to notice because family-level behavior often remains plausible. A user asking for Claude usually receives something Claude-like; the failure is at the version-contract layer, where regression stability, safety-boundary assumptions, and benchmark comparability are silently weakened.

R6 (`api.axxx.com`) — single-event capture of cross-family substitution. R6 accumulated 6 probes and 1 violation (16.7%) — below our $n \geq 5$ evidence threshold. R6 is nonetheless retained as a body-text case because that single violation supplies an internally consistent evidence chain for the cross-family impersonation pattern, ideal for illustrating how the methodology captures such substitutions.

Of 6 probes claiming `claude-opus-4-6`, 3 returned null verdicts (empty or parse-failure responses), 2 returned `clean_match` against genuine Claude Opus 4.6, and **1 returned `plain_mismatch` with the back-end identified as OpenAI GPT-5.4**. The four detection layers reported, for that single event: (1) surface family weights `openai` = 0.96; (2) behavior column likewise OpenAI-leading; (3)

sub-model deterministic match matched GPT-series cutoff/capability/refusal patterns; and (4) response text contained marketing tone and disclaimer patterns idiosyncratic to OpenAI GPT models. **Four mutually independent evidentiary signals all pointed to GPT-5.4; the verdict layer accordingly emitted `plain_mismatch` at high confidence.**

R6's pattern represents a third typology distinct from R1's (persistent dynamic) and R2's (systematic generalization): **incident-style substitution**. While the expected harm to any given downstream user is lower than persistent substitution, its unpredictability defeats any defense based on long-window averaging — the user may receive a non-claimed model precisely on the call that constitutes their commercially critical task.

This case also explains why endpoint-level percentages alone understate the user-facing problem. If a buyer performs ten casual tests, all ten may appear clean; the eleventh call, embedded inside a production workflow, can still be routed to a different family. The measurement question is therefore not only "what fraction of endpoints are persistently substituting?" but also "can the audit system preserve enough per-call evidence to make an isolated substitution event interpretable after the fact?" R6 answers the latter question affirmatively within the limits of black-box observation.

7 Consumer Harm and Policy Implications

7.1 Asymmetric information and quantified harm

Consumer harm in the resale market does not stem from a single "quality below claim" axis but from compounded asymmetries:

- **B1 Quality asymmetry.** Inter-model output-quality differences on a fixed prompt are typically imperceptible to ordinary users on conversation, copywriting, or translation tasks.
- **B2 Version asymmetry.** Enterprises lock prompt engineering and evaluation criteria to specific dated model versions; version-tag forgery silently invalidates the evaluation criterion.
- **B3 Policy asymmetry.** If the reseller strips the upstream model's safety system prompt, downstream users may mistake the resulting policy relaxation for an inherent feature of the model.
- **B4 Litigation asymmetry.** Reseller–buyer relationships rest on unilateral terms-of-service; resellers are typically domiciled outside the buyer's jurisdiction.

By the conservative standard from Section 5.3 ($n \geq 5$, violation rate $\geq 20\%$), **at least 1.3% of sampled active reseller endpoints meet our persistent-substitution criterion**.

7.2 Why upstream providers cannot solve this alone

Upstream providers have no effective means to detect substitution at reseller endpoints. Reasons: (1) upstream sees requests from the reseller's API key with valid `model` parameters, cannot distinguish reseller self-use from sanctioned/impersonated downstream use; (2) the upstream's ledger does not record what the downstream user was told the model is; (3) banning resellers contradicts upstream commercial interest; (4) self-claim pollution *also occurs in upstream providers' own first-party traffic* (a user who injects "you are GPT-4" into Anthropic's official API receives a response self-identifying as GPT-4). **Substitution detection in the resale market cannot rely on upstream providers; it must be supplied by independent third-party measurement infrastructure** [1].

7.3 Three implementable disclosure-duty proposals

D1 — Routing transparency. Reseller endpoints should disclose the actually routed upstream model version in each response's metadata (HTTP header or non-semantic JSON field). Buyers can then verify alignment between advertised model and actual routing. Because voluntary adoption would require resellers to give up substitution opacity, D1 must be mandated by regulation or upstream contract.

D2 — Policy-injection disclosure. Resellers that inject system prompts at the upstream layer must disclose this injection's existence and behavioral impact in their public terms. This is a direct application of existing consumer-protection law's material-disclosure principles.

D3 — Third-party verification badge. Reseller endpoints should permit standardized probing by independent verification organizations and expose the results as a machine-readable badge.

8 Discussion and Limitations

8.1 Sample composition: a consumer-driven endpoint set

The 171 endpoints in this study were not obtained by random sampling from a reseller directory. They are the set of endpoints submitted to our verification service, over a 14-day window, by **retail**

customers of resale endpoints — i.e., by consumers who had paid actual currency to purchase access from a given endpoint and who had reason to doubt the authenticity of the model they had purchased. The population we are entitled to draw conclusions about is therefore "endpoints that some retail customer had reason to verify," not "all resale endpoints in existence" [9]. Our prevalence figures (1.3% strict, 9.9% relaxed) describe *the rate at which consumer-suspected endpoints turn out to actually exhibit substitution* — itself a directly meaningful policy quantity, because it characterizes the endpoints that real consumers paid for and developed real concerns about.

8.2 Detection ceiling

Layer ④ exhibits a 17% false-negative rate (1/6) under injection-mode prompt attacks. The mechanism: when the system prompt explicitly instructs the model to refuse on category X, the model's refusal boundary is pushed up to approximate that of a higher-tier sibling. Closing this ceiling requires either (a) channels resistant to system-prompt rewriting (multi-turn consistency, long-context recall accuracy) or (b) adversarial training across the system-prompt perturbation space.

8.3 Synthetic-test sample size

The 95% Wilson CI of $[74.2\%, 99.0\%]$ on $n{=}18 + 6$ is statistically insufficient to anchor the 94.4% point estimate. The synthetic result should therefore be read as an engineering validation of the detector's directionality rather than as a final performance estimate. Future work should validate the synthetic point estimate at higher n and across a broader set of system-prompt perturbations.

9 Conclusion

By means of a four-layer independent-fingerprint detection system applied to 171 LLM API resale endpoints over 14 days, this paper has quantified for the first time the prevalence and distribution of model substitution in a market that exists in a regulatory gray zone. Under conservative criteria, **at least 1.3% of sampled active endpoints meet our persistent-substitution criterion**; substitutions sort into five categories — cross-family impersonation, silent downgrade, silent upgrade, version-tag forgery, and behavior injection — each with a documented anchor case. The paper analyzes the multi-axis asymmetric-information harm imposed on downstream buyers from a consumer-protection vantage and proposes three immediately implementable disclosure duties.

Future work proceeds along three axes: (1) random-sampling market measurement to eliminate selection bias; (2) closing the detection ceiling by introducing channels resistant to injection-mode prompt attacks; (3) policy operationalization advancing the D1/D2/D3 disclosure duties through industry self-regulation and regulator dialogue.

The LLM API resale market is a gray economy born of access policy in a present regulatory vacuum; the risks borne by its consumers have, for the first time, been systematically and quantitatively exposed. The authors hope this measurement provides a concrete factual basis for market transparency, consumer protection, and forthcoming regulatory dialogue.

References

[1] George A. Akerlof. The market for "lemons": Quality uncertainty and the market mechanism. *The Quarterly Journal of Economics*, 84(3):488–500, 1970.

[2] Anthropic. Anthropic usage policies, 2026. Accessed 2026-04-26.

[3] Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction, 2024.

[4] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know, 2022.

[5] Hanzhi Liu, Chaofan Shou, Hongbo Wen, Yanju Chen, Ryan Jingyang Fang, and Yu Feng. Your agent is mine: Measuring malicious intermediary attacks on the llm supply chain, 2026.

[6] Mason McGovern et al. Stylistic fingerprinting of large language models via output-distribution statistics, 2024.

[7] OpenAI. OpenAI usage policies, 2026. Accessed 2026-04-26.

[8] Dario Pasquini, Eugenia Cretu, and Carmela Troncoso. LLMmap: Fingerprinting for large language models, 2024.

[9] Vern Paxson. Strategies for sound internet measurement. In *Proceedings of the 4th ACM SIGCOMM Internet Measurement Conference (IMC '04)*, pages 263–271. ACM, 2004.

A Ethical Considerations

Disclosure scope. Endpoints are referred to in the body by codenames (R1–R15) with partial domain masking. The full codename–domain mapping is preserved in the authors' internal records; bona-fide researchers may request access via the open-source repository at https://github.com/Bazaarlinkorg/LLMprobe-engine. The codename + masking convention is intended to avoid asymmetric competitive information being injected into the market by this paper, and to preserve rebuttal and re-test space for the named endpoints.

Data minimization. The measurement database stores only response content, token counts, and endpoint identifiers necessary for probe analysis. No personally identifying information, no end-user data of the studied endpoints, and no operations exceeding the publicly advertised access scope of the endpoints were collected.

Measurement constraints. All probe traffic was issued via legitimately purchased API keys, of the same class any retail customer could issue. The study performed no denial-of-service, no billing-bypass, no unauthorized access, and no action exceeding endpoints' publicly advertised access scope. Endpoints were also receiving research-purpose stress-test traffic during the measurement window; a fraction of probe responses flagged as substitution may have arisen from defensive responses to that adversarial traffic. This "test-reflexivity" concern should be weighed especially carefully in the low-evidence cohort ($n<5$).

Dual-use risk and mitigations. The detection methods described here can be used by resellers to evade detection. The authors judge the consumer-protection benefit of disclosure to substantially outweigh this risk. Mitigations: (1) the full codename–domain mapping is held in private records; (2) the open-source implementation is released under copyleft (AGPL-3.0) to maximize independent verification capacity; (3) method exposition focuses on detection techniques usable by consumers for self-protection rather than on attack vectors.

IRB statement. The study does not involve human subjects, human-derived biological material, or personally identifying data. The objects of measurement are publicly marketed Internet-accessible HTTP API services. Consequently, no Institutional Review Board approval was sought or required under 45 CFR 46 (Common Rule, U.S.) or comparable jurisdictional analogues. The authors performed a self-assessment against the ACM Code of Ethics: (1.2 *avoid harm*) no DoS, no resource exhaustion, no billing-bypass, no unauthorized access; (2.3 *respect for rules*) all measurements under endpoints' published terms; (2.6 *competence*) open-source release enables independent verification; (3.1 *public good*) disclosure-duty proposals are framed for retail-consumer benefit.

B Endpoint Codename Statistics

In this arXiv version, the full codename–domain mapping is withheld. The codename–statistics correspondence required to interpret Section 6 is given in Table 1. Bona-fide researchers may request the full mapping via the open-source repository.

C Open-Source Implementation Correspondence

The methodology is released under AGPL-3.0 at https://github.com/Bazaarlinkorg/LLMprobe-engine as the package `@bazaarlink/probe-engine` (latest release v0.7.1 as of this manuscript).

The open-source release ships with offline baseline snapshots for the 22 models reported in Section 5.1, allowing reproduction without dependency on the authors' SaaS API.

Layer	Paper section	Source file
Surface fingerprint	Section 3.3	`src/fingerprint-extractor.ts`
Behavior column	Section 3.4	`src/fingerprint-features-v2.ts`
Sub-model deterministic match	Section 3.5	`src/sub-model-classifier-v3.ts`
Behavioral-vector extension	Section 3.6	`src/sub-model-classifier-v3e.ts` + `-v3f.ts`
Verdict fusion	Section 3.7	`src/identity-verdict.ts`